

# Revisiting structure-exploiting optimal power flow methods

François Pacaud  
*Assistant professor @ Mines Paris-PSL*

**LANS seminar**  
May 24, 2023

# Who I am?

## Short bio

→ numerical optimizer by heart

- Former LANS postdoc, supervised by Mihai Anitescu
- Now assistant professor at Mines Paris-PSL

Today, presenting the work I did during my postdoc @ LANS (2020-2022)

## Research question


Can we solve large-scale nonlinear optimization problems on GPUs?

Fellow collaborators :

- Mihai Anitescu
- Adrian Maldonado
- Michel Schanen
- Sungho Shin

# Motivation: solving optimal power flow problems on GPU architectures

*Our research is funded by the Exascale Computing Project (ECP)*



## Observation

Handling **unstructured sparsity** on **SIMD architectures** is non trivial

Physical model  
**unstructured**



Upcoming hardware  
**GPU centric (SIMD)**



## Why GPUs are hard for optimizers?

## Observation


- GPUs are SIMD architecture (single instruction, multiple data)
- Excellent for *dense* and *batch* operations

On their hand, numerical optimization depends on two key routines

1. **Derivatives:** Evaluate derivatives using *Automatic Differentiation*
2. **Linear solve:** Solve KKT system to compute descent direction  $d_k$  as

$$(\nabla_{xx}^2 \ell_k) d_k = -\nabla_x \ell_k$$

where  $(\nabla_{xx}^2 \ell_k)$  is *sparse symmetric indefinite*



- No good *sparse symmetric indefinite* solver on GPU
- Usual workarounds:
  1. Use decomposition algorithms (ADMM, [Kim et al., 2021])
  2. Use iterative solver (CG-based) [Cao et al., 2016, Schubiger et al., 2020]

(from [Tasseff et al., 2019])

# Our solution: densification

*Intuition: dense is easy on the GPU*

Idea: Exploit the available degrees of freedom

Densify the problem using the *reduced Hessian*

$$\hat{H}_{uu} = Z^\top H Z$$

- Approach widely used during the 1980s/1990s
  - Summarized in [Fletcher, 1987, Section 12.5]: "Nonlinear elimination and feasible direction methods"
  - Also known as "Projected Hessian"  
[Nocedal and Overton, 1985, Gurwitz and Overton, 1989]
  - The reduced Hessian  $\hat{H}_{uu}$  is often approximated [Biegler et al., 1995]
- The optimization community moved away from this technique in the 2000s:
  - "Many degrees of freedom" approaches [Poku et al., 2004]
  - Efficient resolution with interior-point combined with generic indefinite sparse linear solver (HSL [Duff and Reid, 1983], Pardiso [Schenk and Gärtner, 2004])
  - Lead to the development of mature NLP solvers (Ipopt [Wächter and Biegler, 2006], Knitro [Waltz et al., 2006])



MadNLP [Shin et al., 2020]

- Port of Ipopt in Julia
- Filter line-search interior-point method
- Open-source:  
<https://github.com/MadNLP/MadNLP.jl>

- **Derivatives:**
  - Custom automatic-differentiation backend for OPF: ExaPF.jl
  - Derivatives evaluated in parallel **on the GPU**
- **Linear solver:** We compare two linear solvers for the KKT system
  1. *The reference:* HSL ma27 running **on the CPU**
  2. *The contender:* our reduction algorithm, using cusolver to factorize the reduced matrix with dense Cholesky **on the GPU**

# Outline

Optimal power flow (OPF)

Security-Constrained Optimal power flow (SC-OPF)

Transient stability constraint OPF (TS-OPF)

# Expliciting the physical constraints in the optimization problem




Figure: Nonlinear power flow  
(from [Hiskens and Davy, 2001])

Most real-life nonlinear problems encompasses a set of physical constraints

$$g(\mathbf{x}, \mathbf{u}) = 0$$

with  $\mathbf{x}$  a state and  $\mathbf{u}$  a control

| Domain                       | $g$               |
|------------------------------|-------------------|
| Optimal control              | Dynamics          |
| PDE-constrained optimization | PDE               |
| <b>Optimal power flow</b>    | <b>Power flow</b> |

## Physically-constrained optimization problem

$$\min_{\mathbf{x}, \mathbf{u}} f(\mathbf{x}, \mathbf{u})$$

$$\text{s.t. } g(\mathbf{x}, \mathbf{u}) = 0, \quad h(\mathbf{x}, \mathbf{u}) \leq 0$$

Well-known method [Cervantes et al., 2000, Biros and Ghattas, 2005]



## Interior-point in a nutshell

### Notations

- $W$ : Hessian of Lagrangian
- $G$ : Jacobian of equality constraints (power flow)
- $A$ : Jacobian of inequalities (operational constraints)

The interior-point method (IPM) rewrites the problem in the standard form:



$$\begin{aligned} \min_{\mathbf{x}, \mathbf{u}, \mathbf{s}} \quad & f(\mathbf{x}, \mathbf{u}) \\ \text{s.t.} \quad & \mathbf{g}(\mathbf{x}, \mathbf{u}) = 0, \quad \mathbf{h}(\mathbf{x}, \mathbf{u}) + \mathbf{s} = 0, \quad \mathbf{s} \geq 0 \end{aligned}$$

At each iteration, it solves the **augmented KKT linear system**

$$\begin{bmatrix} W + \Sigma_p & 0 & G^T & A^T \\ 0 & \Sigma_s & 0 & I \\ G & 0 & 0 & 0 \\ A & I & 0 & 0 \end{bmatrix} \begin{bmatrix} \mathbf{p}_v \\ \mathbf{p}_s \\ \mathbf{p}_\lambda \\ \mathbf{p}_y \end{bmatrix} = \begin{bmatrix} \mathbf{r}_1 \\ \mathbf{r}_2 \\ \mathbf{r}_3 \\ \mathbf{r}_4 \end{bmatrix}$$

- in olive, blocks associated to the inequality constraints
- in violet, blocks associated to the equality constraints

## Condense step: we remove the inequality constraints



We remove the **blocks** associated to the inequality constraints by taking the Schur-complement

### Condensed KKT

We define the *condensed KKT matrix* as

$$K := W + \mathbf{A}^\top \Sigma_s \mathbf{A}$$



The augmented KKT system is equivalent to

$$\begin{bmatrix} K + \Sigma_p & G^\top \\ G & 0 \end{bmatrix} \begin{bmatrix} \mathbf{p}_p \\ \mathbf{p}_\lambda \end{bmatrix} = \begin{bmatrix} \mathbf{r}_1 + \mathbf{A}^\top (\Sigma_s \mathbf{r}_4 + \mathbf{r}_2) \\ \mathbf{r}_3 \end{bmatrix}$$

N.B.: This step is usually discarded because of additional fill-in in left-hand-side matrix, but here our goal is to **densify** the KKT system

## Reduce step: we remove the equality constraints

**Idea:** exploit the structure of the power flow equations  $g(\mathbf{x}, \mathbf{u}) = 0$



Expanding the structure of the condensed KKT system:

$$\begin{bmatrix} K_{xx} + \Sigma_x & K_{xu} & G_x^T \\ K_{ux} & K_{uu} + \Sigma_u & G_u^T \\ G_x & G_u & 0 \end{bmatrix} \begin{bmatrix} \mathbf{p}_x \\ \mathbf{p}_u \\ \mathbf{p}_\lambda \end{bmatrix} = \begin{bmatrix} \hat{\mathbf{r}}_1 \\ \hat{\mathbf{r}}_2 \\ \hat{\mathbf{r}}_3 \end{bmatrix}$$

### Reduced KKT

If the Jacobian  $G_x$  is invertible, the reduced Hessian is defined as

$$\hat{K}_{uu} := Z^T K Z \quad \text{with} \quad Z := \begin{bmatrix} -G_x^{-1} G_u \\ I \end{bmatrix}$$

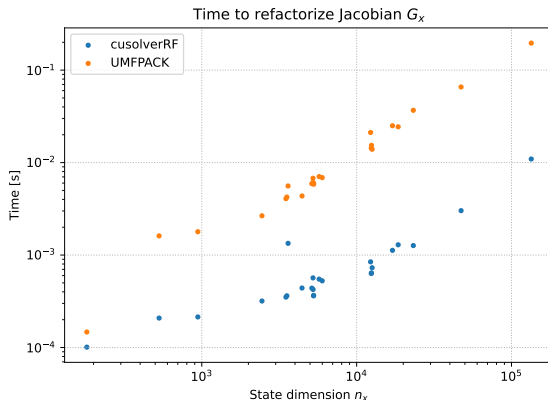
The condensed KKT system is equivalent to

$$\hat{K}_{uu} \mathbf{p}_u = \hat{\mathbf{r}}_2 - G_u^T G_x^{-T} \hat{\mathbf{r}}_1 - (K_{ux} - G_u^T G_x^{-T} K_{xx}) G_x^{-1} \hat{\mathbf{r}}_3$$

- Assembling  $\hat{K}_{uu}$  can proceed in parallel once the sparse Jacobian  $G_x$  factorized
- The matrix  $\hat{K}_{uu}$ , **dense**, can be factorized efficiently on the GPU

## Step 1: factorizing $G_x$

**Message 1:** cusolverRF is efficient to refactorize the matrix  $G_x$  on CUDA GPU



**Figure:** Time to refactorize the Jacobian  $G_x$  against the dimension of the network.

## Step 2: parallel reduction algorithm (1)

We suppose given the condensed KKT matrix  $K = W + A^\top \Sigma_s A$

$$\hat{K}_{uu} = \begin{bmatrix} -G_x^{-1} G_u \\ I \end{bmatrix}^\top \begin{bmatrix} K_{xx} & K_{xu} \\ K_{ux} & K_{uu} \end{bmatrix} \begin{bmatrix} -G_x^{-1} G_u \\ I \end{bmatrix}$$

We should avoid allocating the sensitivity matrix  $S = -G_x^{-1} G_u$  (dense, size  $n_x \times n_u$ )!

→ Instead, use batched HessMat product  $\hat{K}_{uu} V$

HessMat kernel: batch adjoint-adjoint reduction

Input: LU factorization, such that  $PG_x Q = L U$  (2 SpMM, 2 SpSM)

For every RHS  $V \in \mathbb{R}^{n_u \times N}$

1. Solve  $Z = G_x^{-1} (G_u V)$  (3 SpMM, 2 SpSM)

2. Evaluate  $\begin{bmatrix} \Psi \\ H_u \end{bmatrix} = \begin{bmatrix} K_{xx} & K_{xu} \\ K_{ux} & K_{uu} \end{bmatrix} \begin{bmatrix} Z \\ V \end{bmatrix}$  (1 SpMM)


3. Solve  $H_x = G_x^{-\top} \Psi$  (2 SpMM, 2 SpSM)

4. Output  $\hat{K}_{uu} V = H_u - G_u H_x$  (1 SpMM)

- $G_x$  first factorized on the CPU with KLU, then refactorized entirely on the GPU with `cusolverRF` (fast)
- $\text{div}(n_u, N) + 1$  HessMat products required to get full  $\hat{K}_{uu}$

## Step 2: parallel reduction algorithm (2)

**Message 2:** 7 seconds to evaluate  $\hat{K}_{uu}$  for the largest instance (ACTIVSg70k)



**Figure:** Time to evaluate  $\hat{K}_{uu}$  against the number of batch on the GPU, for various networks.

# Benchmark 1: Porting automatic differentiation on the GPU

## Setting

- **Knitro+ma27**
  - **Derivatives:** MATPOWER (matlab)
  - **Linear solver:** ma27
- **MadNLP+ma27**
  - **Derivatives:** GPU-accelerated AD
  - **Linear solver:** ma27

|            | Knitro+ma27 |              | MadNLP+ma27 |             |
|------------|-------------|--------------|-------------|-------------|
| Case       | #it         | Time (s)     | #it         | Time (s)    |
| 9241pegase | 102         | <b>31.7</b>  | 69          | <b>10.7</b> |
| ACTIVSg25k | 47          | <b>36.1</b>  | 86          | <b>24.7</b> |
| ACTIVSg70k | 101         | <b>242.0</b> | 90          | <b>89.8</b> |
| 9591goc    | 37          | <b>22.5</b>  | 43          | <b>11.7</b> |
| 10480goc   | 40          | <b>25.7</b>  | 50          | <b>14.0</b> |
| 19402goc   | 45          | <b>66.5</b>  | 47          | <b>30.8</b> |

## Benchmark 2: Porting the linear solver on the GPU

### Setting

- **MadNLP+ma27**
  - **Derivatives:** GPU-accelerated AD
  - **Linear solver:** ma27
- **MadNLP+reduced KKT**
  - **Derivatives:** GPU-accelerated AD
  - **Linear solver:** reduction on GPU

|  |      | <i>The reference</i><br><b>MadNLP+ma27</b> |             |          | <i>The contender</i><br><b>MadNLP+reduced KKT</b> |              |           |               |
|--|------|--|-------------|----------|---|--------------|-----------|---------------|
| Case   | DOF  | #it  | Time (s)    | ma27 (s) | #it   | Time (s)     | Chol. (s) | Reduction (s) |
| <b>Problems with many degrees of freedom</b> |      |  |             |          |   |              |           |               |
| 9241pegase                                   | 0.14 | 69   | <b>10.7</b> | 6.1      | 69  | <b>23.7</b>  | 1.2       | 16.2          |
| ACTIVSg25k                                   | 0.10 | 86   | <b>24.7</b> | 16.9     | 86  | <b>85.0</b>  | 4.3       | 68.1          |
| ACTIVSg70k                                   | 0.08 | 90   | <b>89.8</b> | 65.7     | 85  | <b>658.2</b> | 21.5      | 606.5         |
| <b>Problems with few degrees of freedom</b>  |      |  |             |          |   |              |           |               |
| 9591goc                                      | 0.02 | 43   | <b>11.7</b> | 10.4     | 43  | <b>7.7</b>   | 2.1       | 1.6           |
| 10480goc                                     | 0.03 | 50   | <b>14.0</b> | 12.0     | 50  | <b>11.5</b>  | 3.9       | 3.3           |
| 19402goc                                     | 0.02 | 47   | <b>30.8</b> | 26.8     | 47  | <b>19.5</b>  | 4.9       | 7.2           |

**Table:** Comparing ma27 with reduced KKT linear solver. DOF is the % of degrees of freedom.



# When is reduced better than full-space?

## Observation

The smaller the number of degrees of freedom  $n_u$ , the more efficient is the reduction of the KKT system




Figure: Illustrating the breakeven point

# Summary for the classical OPF problem

## Where we are at?

- **Derivatives:**
  - $\approx 2$  times speed-up compared to MATPOWER just by evaluating the derivatives on GPU
- **Linear solver:**
  - Reduced KKT is a new structure-exploiting linear solver
  - Implemented entirely on the GPU (without memory transfer / the host)
  - Tractable in the large-scale regime, and beat state-of-the-art if # DOF is small

# Outline

Optimal power flow (OPF)

Security-Constrained Optimal power flow (SC-OPF)

Transient stability constraint OPF (TS-OPF)

### SCOPF

Add  $N$  contingency scenarios (line tripping) to the base case OPF

In **preventive mode**, the SCOPF formulates as

$$\min_{x_0, x_1, \dots, x_N, u} f(x_0, u) \quad \text{subject to} \quad \begin{cases} g(x_0, u) = 0 \\ h(x_0, u) \leq 0 \\ g(x_c, u) = 0 \quad \forall c = 1, \dots, N \\ h(x_c, u) \leq 0 \quad \forall c = 1, \dots, N \end{cases}$$

Conservative formulation: the control  $u$  (=power generations) is shared across all contingencies

### Observations

- The derivatives can be evaluated in batch on the GPU
- The associated KKT system has a block arrowhead structure we can exploit in the reduced KKT solver

We follow the same procedure as before: *condense* and *reduce*

## Fact 1: the condensed matrix has a block-arrowhead structure

The Hessian and the Jacobian of the SCOPF problems have all a *block-arrowhead structure*

$$W = \begin{bmatrix} W_{x_1 x_1} & & & W_{x_1 u} \\ & \ddots & & \vdots \\ & & W_{x_N x_N} & W_{x_N u} \\ W_{ux_1} & \dots & W_{ux_N} & W_{uu} \end{bmatrix}, \quad A = \begin{bmatrix} A_{x_1}^1 & & & A_u^1 \\ & \ddots & & \vdots \\ & & A_{x_N}^N & A_u^N \end{bmatrix}$$




Figure: Condensed matrix  $K$

### Proposition

The condensed matrix  $K = W + A^T \Sigma_s A$  inherits the block-arrowhead structure of the Hessian  $W$  and the Jacobian  $H$ , with

$$K = \begin{bmatrix} K_{x_1 x_1} & & & K_{x_1 u} \\ & \ddots & & \vdots \\ & & K_{x_N x_N} & K_{x_N u} \\ K_{ux_1} & \dots & K_{ux_N} & K_{uu} \end{bmatrix}$$

## Fact 2: we can reduce the KKT system by blocks

**Idea:** Exploit the block-arrowhead structure in the reduction

### Proposition

If the Jacobian  $G_x^0, \dots, G_x^N$  are invertible, the reduced Hessian is defined as

$$\hat{K}_{uu} := Z^\top K Z \quad \text{with} \quad Z := \begin{bmatrix} -(G_x^0)^{-1} G_u^0 \\ \vdots \\ -(G_x^N)^{-1} G_u^N \\ I \end{bmatrix}$$

The KKT system is equivalent to

$$\hat{K}_{uu} p_u = -\hat{r}_2 + \sum_{i=1}^N \left[ (G_u^i)^\top (G_x^i)^{-\top} \hat{r}_1^i + [K_{ux_i} - (G_u^i)^\top (G_x^i)^{-\top} K_{x_i x_i}] (G_x^i)^{-1} \hat{r}_3^i \right]$$

- The Jacobians  $G_x^0, \dots, G_x^N$  can be factorized in parallel on the GPU using `cusolverRF`
- The reduced Hessian  $\hat{K}_{uu}$  is dense, with dimension  $n_u \times n_u$

# Numerical results on a single GPU (NVIDIA V100)

Setting: same as before

- **MadNLP+ma27**
  - **Derivatives:** GPU-accelerated AD
  - **Linear solver:** ma27
- **MadNLP+reduced KKT**
  - **Derivatives:** GPU-accelerated AD
  - **Linear solver:** reduction on GPU

| #bus | N  | MadNLP+ma27 |               |        |          | MadNLP+reduced KKT |               |        |               |
|------|----|-------------|---------------|--------|----------|--------------------|---------------|--------|---------------|
|      |    | #it         | Tot (s)       | AD (s) | ma27 (s) | #it                | Tot (s)       | AD (s) | reduction (s) |
| 1354 | 8  | 61          | <b>10.8</b>   | 0.9    | 9.9      | 61                 | <b>7.9</b>    | 0.9    | 7.0           |
| 1354 | 16 | 54          | <b>26.2</b>   | 1.1    | 25.1     | 54                 | <b>13.3</b>   | 1.2    | 12.1          |
| 1354 | 32 | 253         | <b>1302.0</b> | 9.0    | 1293.0   | 233                | <b>172.2</b>  | 9.0    | 163.2         |
| 1354 | 64 | 135         | <b>411.7</b>  | 8.6    | 403.1    | 236                | <b>357.3</b>  | 14.5   | 342.8         |
| 9241 | 8  | 190         | <b>1400.5</b> | 31.7   | 1368.8   | 187                | <b>1017.0</b> | 30.5   | 986.5         |
| 9241 | 16 | 121         | <b>3947.0</b> | 38.1   | 3908.9   | 123                | <b>1091.4</b> | 34.4   | 1067.0        |

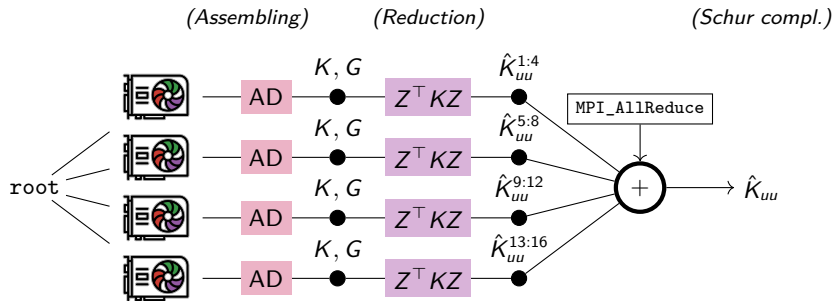
**Table:** Comparing the performance of the KKT linear solvers

## Extension to multiple GPUs

### Observation

We can gain additional speed-up by dispatching the solution on multiple GPUs

- ✓ Porting OPF to the exascale regime
- ✓ The approach is equivalent to the Schur-decomposition implemented previously in PIPS-NLP





## Performance of automatic differentiation on multiple GPUs




Figure: Time spent to evaluate the derivatives of case1354pegase with automatic differentiation

# Performance of parallel reduced KKT solver on multiple GPUs





Figure: Time spent to solve the KKT system for case1354pegase

## Solving SCOPF on multiple GPUs



**Figure:** Time to solve the block-structured OPF problem case118 as a function of the number of scenarios  $N$

## Where we are at?

- **Derivatives:**
  - Very fast evaluation of the derivatives in batch on the GPU
- **Linear solver:**
  - Reduced KKT offers very good performance for SCOPF, as #DOF remains small (control  $u$  shared across contingencies)
  - Reduction can run in parallel on multiple GPUs
- **Generic observation:**
  - ✓ GPUs are well adapted to solve large-scale block-NLP problems with a shared structure (here: topology of the network is fixed +/- one line contingency)

# Outline

Optimal power flow (OPF)

Security-Constrained Optimal power flow (SC-OPF)

Transient stability constraint OPF (TS-OPF)

## Our next-step: transient-stability constraint problems

Extension to transient stability constraint OPF is under way

- Formulate as a NLP with DAE constraints
- Reduced-space approach for TS-OPF already studied in [Jiang and Geng, 2010]

### Observation

The Jacobian inherits the dynamics structure of the problem:

$$G_x = \begin{bmatrix} G_{M,x_1} & & & \cdots & \\ G_{S,x_2} & G_{M,x_2} & & \cdots & \\ & G_{S,x_3} & G_{M,x_3} & \cdots & \\ \vdots & \cdots & \cdots & \ddots & \vdots \\ & & & G_{S,x_T} & G_{M,x_T} \end{bmatrix}$$

It can also be re-interpreted as an optimal control problem

- [Wright, 1993, Cervantes et al., 2000]
- Parallel linear solver for dynamic problems already exist [Wright, 1990, Wright, 1991]

## References I



Biegler, L. T., Nocedal, J., and Schmid, C. (1995).  
A reduced Hessian method for large-scale constrained optimization.  
*SIAM Journal on Optimization*, 5(2):314–347.



Biros, G. and Ghattas, O. (2005).  
Parallel Lagrange–Newton–Krylov–Schur Methods for PDE-Constrained Optimization. Part I: The Krylov–Schur Solver.  
*SIAM Journal on Scientific Computing*, 27(2):687–713.



Cao, Y., Seth, A., and Laird, C. D. (2016).  
An augmented Lagrangian interior-point approach for large-scale NLP problems on graphics processing units.  
*Computers & Chemical Engineering*, 85:76–83.



Cervantes, A. M., Wächter, A., Tütüncü, R. H., and Biegler, L. T. (2000).  
A reduced space interior point strategy for optimization of differential algebraic systems.  
*Computers & Chemical Engineering*, 24(1):39–51.



Duff, I. S. and Reid, J. K. (1983).  
The multifrontal solution of indefinite sparse symmetric linear.  
*ACM Transactions on Mathematical Software (TOMS)*, 9(3):302–325.

## References II



Fletcher, R. (1987).  
*Practical methods of optimization*.  
John Wiley & Sons.



Gurwitz, C. B. and Overton, M. L. (1989).  
Sequential quadratic programming methods based on approximating a  
projected Hessian matrix.  
*SIAM Journal on Scientific and Statistical Computing*, 10(4):631–653.



Hiskens, I. A. and Davy, R. J. (2001).  
Exploring the power flow solution space boundary.  
*IEEE transactions on power systems*, 16(3):389–395.



Jiang, Q. and Geng, G. (2010).  
A reduced-space interior point method for transient stability constrained  
optimal power flow.  
*IEEE Transactions on Power Systems*, 25(3):1232–1240.



Kim, Y., Pacaud, F., Kim, K., and Anitescu, M. (2021).  
Leveraging gpu batching for scalable nonlinear programming through  
massive lagrangian decomposition.  
*arXiv preprint arXiv:2106.14995*.



## References III



Nocedal, J. and Overton, M. L. (1985).

Projected Hessian updating algorithms for nonlinearly constrained optimization.

*SIAM Journal on Numerical Analysis*, 22(5):821–850.



Poku, M. Y. B., Biegler, L. T., and Kelly, J. D. (2004).

Nonlinear optimization with many degrees of freedom in process engineering.

*Industrial & engineering chemistry research*, 43(21):6803–6812.



Schenk, O. and Gärtner, K. (2004).

Solving unsymmetric sparse systems of linear equations with pardiso.

*Future Generation Computer Systems*, 20(3):475–487.



Schubiger, M., Banjac, G., and Lygeros, J. (2020).

GPU acceleration of ADMM for large-scale quadratic programming.

*Journal of Parallel and Distributed Computing*, 144:55–67.



Shin, S., Coffrin, C., Sundar, K., and Zavala, V. M. (2020).

Graph-based modeling and decomposition of energy infrastructures.

*arXiv preprint arXiv:2010.02404*.

## References IV



Tasseff, B., Coffrin, C., Wächter, A., and Laird, C. (2019).  
Exploring benefits of linear solver parallelism on modern nonlinear  
optimization applications.  
*arXiv preprint arXiv:1909.08104*.



Wächter, A. and Biegler, L. T. (2006).  
On the implementation of an interior-point filter line-search algorithm for  
large-scale nonlinear programming.  
*Mathematical Programming*, 106(1):25–57.



Waltz, R. A., Morales, J. L., Nocedal, J., and Orban, D. (2006).  
An interior algorithm for nonlinear optimization that combines line search  
and trust region steps.  
*Mathematical Programming*, 107(3):391–408.



Wright, S. J. (1990).  
Solution of discrete-time optimal control problems on parallel computers.  
*Parallel Computing*, 16(2-3):221–237.



Wright, S. J. (1991).  
Partitioned dynamic programming for optimal control.  
*SIAM Journal on optimization*, 1(4):620–642.

## References V



Wright, S. J. (1993).

Interior point methods for optimal control of discrete time systems.

*Journal of Optimization Theory and Applications*, 77(1):161–187.