

# GPU-accelerated optimal control

*A nonlinear programming point-of-view*

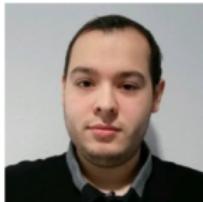
François Pacaud  
[frapac.github.io](https://frapac.github.io)

Centre Automatique et Systèmes, Mines Paris - PSL

Guest lecture  
**Georgia Institute of Technology**  
October 17th, 2025

# Who are we?

<https://madsuite.org/>



- ☞ Alexis Montoison @ ANL
- ☞ Sungho Shin @ MIT
- ☞ Mihai Anitescu @ ANL

Outline: today we talk about GPU-accelerated optimal control

- ☞ We leverage the GPU-accelerated optimization solver MadNLP to solve large-scale optimal control on the GPU

## Motivation for today's lecture

- » New generation of optimization solvers is under way



Figure: Go check Sungho's talk at ScaleOpt!

- » Massive investments in robotics

### Nvidia Bets Big on Robots

The chipmaker, which has led a rally in artificial intelligence stocks, laid out a vision for dominating so-called physical AI. Investors appeared impressed.



By Andrew Ross Sorkin, Ravi Mattu, Bernhard Warner, Sarah Kessler, Michael J. de la Merced and Lauren Hirsch

Jan. 7, 2025



Figure: NVIDIA Jetson

# Outline

What am I talking about when talking about optimal control?

Primal-dual interior-point method

How to solve the Newton systems on the GPU?

Using MadNLP and ExaModels on the GPU

# Applications of optimal control

Finding optimal trajectory in problems with a dynamic structure

- Trajectory planning
- Real-time optimization with model predictive control (MPC)

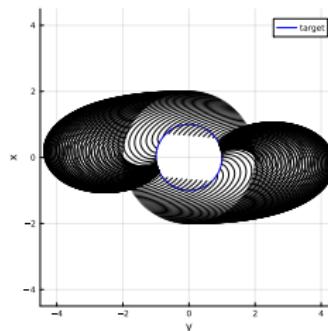


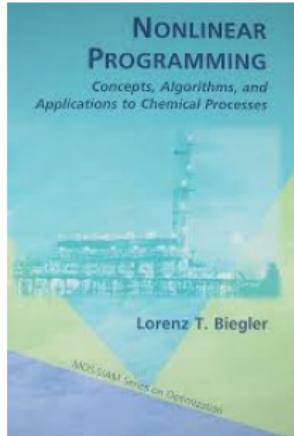
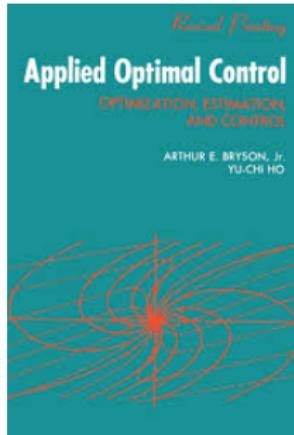
Figure: Trajectory flow for moving a satellite locally close to its orbit

A mature ecosystem already exists!

Classical workflow:

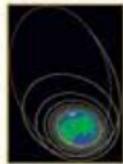
- Modeling with casadi
- Solving with Ipopt (or any structure exploiting solver)

# Useful references



**Practical Methods for  
Optimal Control Using  
Nonlinear Programming**

Third Edition



**John T. Betts**



**Model Predictive Control:  
Theory, Computation, and Design**  
2nd Edition

James B. Rawlings  
David Q. Mayne  
Moritz M. Diehl

# What do I mean by optimal control?

Disclaimer: We write everything in discrete time

At discrete time  $k$ , the *state* of the system is encoded in a vector  $x_k \in \mathbb{R}^{n_x}$ .

After applying a *control*  $u_k \in \mathbb{R}^{n_u}$ , the system evolves between  $k$  and  $k + 1$  as

$$x_{k+1} = g_k(x_k, u_k)$$

Most of the time, the dynamics  $g_k(\cdot)$  discretizes an ordinary differential equations (e.g. using collocation)

The system is subject to the operational constraints:

$$h_k(x_k, u_k) \leq 0$$

## Optimal control problem

Starting from  $\underline{x} \in \mathbb{R}^{n_x}$ , we minimize the functional  $\ell_k(\cdot)$  over an horizon  $N$ :

$$\min_{x,u} \sum_{k=0}^{N-1} \ell_k(x_k, u_k) + \ell_N(x_N)$$

$$\text{s.t. } x_{k+1} = g_k(x_k, u_k), \quad x_0 = \underline{x}$$
$$h_k(x_k, u_k) \leq 0$$

Let  $x := (x_0, \dots, x_N)$  and  $u := (u_0, \dots, u_{N-1})$ .

# Introducing the nonlinear program

Many degrees-of-freedom approach

For  $w := (x, u)$ , we abstract the previous dynamic program as:

$$\begin{aligned} & \min_w f(w) \\ & \text{s.t. } g(w) = 0, \quad h(w) \leq 0. \end{aligned} \tag{NLP}$$

## Nonlinear programming

- Problem is nonlinear nonconvex :  
we are interested in finding only a **local** optimum solution
- Solvable using classical nonlinear solvers:
  - IPM (Ipopt, Knitro, MadNLP)
  - SQP (FilterSQP, SNOPT)
  - Augmented-Lagrangian (LANCELOT, Algencan)

## Karush-Kuhn-Tucker (KKT) conditions

Using a slack variable  $s \geq 0$ , NLP is equivalent to

$$\begin{array}{ll} \min_{w,s} & f(w) \\ \text{s.t. } & g(w) = 0, h(w) + s = 0 \\ & s \geq 0 \end{array} \quad \text{(NLP)}$$

Slack

We introduce the *Lagrangian*:

$$\mathcal{L}(w, y, z) = f(w) + y^\top g(w) + z^\top h(w).$$

### KKT stationary equations

If  $w$  is a *regular* local solution, then there exist dual multipliers  $(y, z)$  satisfying

$$\begin{cases} \nabla f(w) + \nabla g(w)^\top y + \nabla h(w)^\top z = 0 \\ g(w) = 0 \\ h(w) + s = 0 \\ 0 \leq s \perp z \geq 0 \end{cases}$$

↑ Complementarity conditions

- Solving the NLP is equivalent to the solution of a system of *nonsmooth* nonlinear equations

## Regularity conditions: first-order

Denote the active set  $\mathcal{A}(w) = \{i \in [m] \mid h_i(w) = 0\}$  and the active Jacobian

$$J(w) = \begin{bmatrix} \nabla g(w) \\ \nabla h_{\mathcal{A}}(w) \end{bmatrix}$$

### First-order constraint qualification

We say that the point  $w$  is *qualified* if the local geometry of the feasible can be captured by a *linearized* model near  $w$

(in math language, the tangent cone is equal to the set of linearized feasible directions)

- The *Linear Independence Constraint Qualification* (LICQ) holds if the active Jacobian  $J(w)$  is full row-rank
- The *Mangasarian-Fromovitz Constraint Qualification* (MFCQ) holds if  $(y, z) = (0, 0)$  is the unique solution to the linear system

$$\begin{cases} \nabla g(w)^T y + \nabla h(w)^T z = 0 \\ z_i \geq 0 \quad \forall i \in \mathcal{A}(w) \end{cases}$$

### Gauvin's theorem (1979)

MFCQ holds if and only if the set of multipliers  $(y, z)$  satisfying KKT is bounded

## Regularity conditions: second-order

Let  $(w, y, z)$  a primal-dual point satisfying KKT. We define the *critical cone*:

$$\begin{aligned}\mathcal{C}(w, y, z) = \{d \in \mathbb{R}^n \mid & \nabla g_i(w)^\top d = 0, i \in [m_e], \\ & \nabla h_i(w)^\top d = 0, i \in \mathcal{A}(w) \text{ with } z_i > 0, \\ & \nabla h_i(w)^\top d \leq 0, i \in \mathcal{A}(w) \text{ with } z_i = 0\}.\end{aligned}$$

*Strict complementarity (SC)* holds if  $z_i > 0$  for all  $i \in \mathcal{A}(w)$ .

### Second-order sufficient condition (SOSC)

We say that  $(w, y, z)$  satisfies SOSC if

$$d^\top \nabla_{ww}^2 \mathcal{L}(w, y, z) d > 0 \quad \forall d \in \mathcal{C}(w, y, z) \setminus \{0\}$$

Under SOSC, the problem is locally convex near  $w$  and the solution is isolated

### Proposition

Suppose LICQ and SC hold. Then SOSC holds if and only if

$$Z^\top \nabla_{ww}^2 \mathcal{L}(w, y, z) Z \succ 0$$

with  $Z$  a basis of the null-space of the active Jacobian  $J(w)$ .

# Are optimal control problems regular?

Well... it depends

## State constraints

- Pure state constraints:  $h_k(x_k) \leq 0$
  - Mixed state constraints:  $h_k(x_k, u_k) \leq 0$
- 
- ☞ Problems with active *state constraints* do not satisfy LICQ  
(not enough degrees of freedom)
  - ☞ Problems with *singular arcs* do not satisfy SOSC  
Example: Goddard rocket's problem

## Observation

Oftentimes, nonlinear solvers are struggling to solve optimal control instances

# Outline

What am I talking about when talking about optimal control?

**Primal-dual interior-point method**

How to solve the Newton systems on the GPU?

Using MadNLP and ExaModels on the GPU

## Interior-point method (IPM)

Rewrite the (nonsmooth) KKT system as a *smooth* nonlinear system

$$F_\mu(x, s; \boxed{y, z}) := \begin{bmatrix} \nabla f(x) + \nabla g(x)^\top \boxed{y} + \nabla h(x)^\top \boxed{z} \\ g(x) \\ h(x) + s \\ S\nu - \mu e \end{bmatrix} = 0$$

$\uparrow$  Homotopy,  $S = \text{diag}(s)$

### Primal-dual interior-point method

Solve  $F_\mu(x, s; y, z, \nu) = 0$  using Newton method while driving  $\mu \rightarrow 0$ .

# Augmented KKT system

The Newton algorithm translates to the solution of a sequence of linear systems

Denote the primal-dual variable by  $v^k := (w^k, s^k, y^k, z^k)$ . At iteration  $k$ ,

1. Compute Newton step  $d^k$  as solution of the linear system

$$\nabla F_\mu(v^k) d^k = -F_\mu(v^k)$$

2. Update the as

$$v^{k+1} = v^k + \alpha^k d^k$$

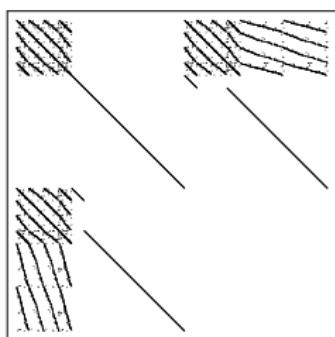


Figure:  $\nabla F_\mu$

## Augmented KKT system

After (slight) reformulation, the Newton step writes as

$$\begin{bmatrix} W & 0 & \nabla g^\top & \nabla h^\top \\ 0 & \Sigma_s & 0 & I \\ \nabla g & 0 & 0 & 0 \\ \nabla h & I & 0 & 0 \end{bmatrix} \begin{bmatrix} d_w \\ d_s \\ d_y \\ d_z \end{bmatrix} = - \begin{bmatrix} r_1 \\ r_2 \\ r_3 \\ r_4 \end{bmatrix}$$

with  $W = \nabla_{xx}^2 L(\cdot)$ ,  $\Sigma_s = S^{-1} \text{diag}(z)$

☞ This is the system solved by default in Ipopt

## Generic case: sparse linear solver for nonlinear programming

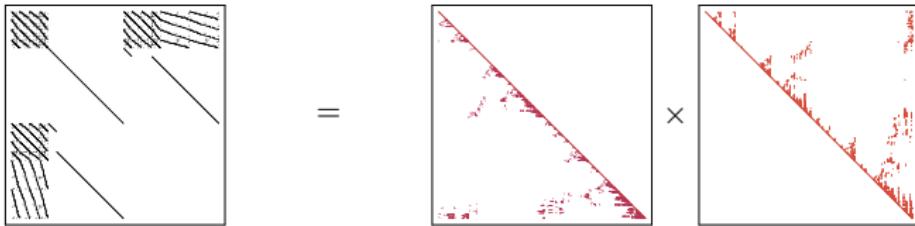


Figure: Matrix factorization using a direct solver

### Direct method for sparse indefinite linear systems

- Ill-conditioning of the KKT system  
*(= iterative solvers are often not practical)*
- Direct solver requires **numerical pivoting** for stability  
*(= difficult to parallelize)*

## Duff-Reid factorization

$$A = P Q L B L^\top Q^\top P^\top$$

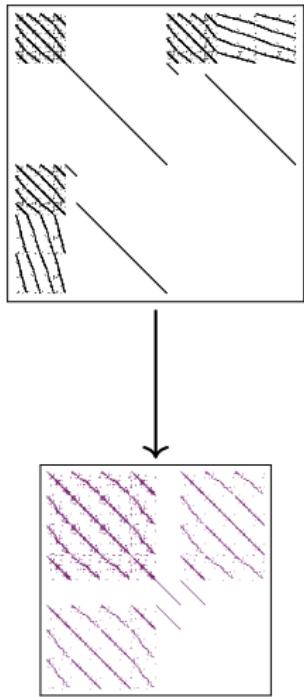
with

- $P$ : fill-in minimization matrix
- $Q$ : additional pivoting for numerical stability
- $L$ : unit lower-triangular matrix
- $B$ : block diagonal matrix with blocks of dimension  $1 \times 1$  or  $2 \times 2$

- ☞ The LBL factorization has become competitive only in the 1990s, using a technique known as matching-based preprocessing
- ☞ The progress in sparse linear solvers has directly benefited to nonlinear optimization solvers such as Ipopt or Knitro
- ☞ Numerical pivoting  $Q$  impairs the parallelism in the algorithm

# Condensed KKT system

We condense the linear system by removing the inequality blocks



## Condensed KKT system

The augmented KKT system condenses to

$$\begin{bmatrix} \textcolor{teal}{K} & \nabla g^\top \\ \nabla g & 0 \end{bmatrix} \begin{bmatrix} d_w \\ d_y \end{bmatrix} = - \begin{bmatrix} r_1 + (\nabla h)^\top (\Sigma_s r_4 + r_2) \\ r_3 \end{bmatrix}$$

with the *condensed matrix*  $K = W + \nabla h^\top \Sigma_s \nabla h$ .

We recover  $(d_s, d_z)$  as

$$d_s = -\Sigma_s^{-1}(r_3 + d_y), \quad d_z = \Sigma_s(\nabla h d_x - r_4) - r_2.$$

- ☞ Additional fill-in
- ☞ Useful when large number of inequality constraints  $m$

# Condensed KKT systems in optimal control

Writing the control as a feedback on the state

## Proposition

For  $K$  positive definite, the solution of the saddle-point linear system

$$\begin{bmatrix} K & G^\top \\ G & 0 \end{bmatrix} \begin{bmatrix} w \\ y \end{bmatrix} = \begin{bmatrix} -c \\ b \end{bmatrix}$$

is the primal-dual solution of the convex QP

$$\min_w \frac{1}{2} w^\top K w + c^\top w \quad \text{s.t.} \quad Gw = b$$

For problem with a dynamic structure, the condensed KKT system yields the LQR

$$\begin{aligned} \min_{d_x, d_u} \quad & \sum_{k=0}^{N-1} \begin{bmatrix} d_{x,k} \\ d_{u,k} \end{bmatrix}^\top \begin{bmatrix} K_{xx}^k & K_{xu}^k \\ K_{ux}^k & K_{uu}^k \end{bmatrix} \begin{bmatrix} d_{x,k} \\ d_{u,k} \end{bmatrix} + \begin{bmatrix} r_{x,k} \\ r_{u,k} \end{bmatrix}^\top \begin{bmatrix} d_{x,k} \\ d_{u,k} \end{bmatrix} \\ \text{s.t.} \quad & d_{x,k+1} = G_{x,k} d_{x,k} + G_{u,k} d_{u,k} \end{aligned}$$

- ☞ No need to use a sparse linear solver
- ☞ Efficient solution with (backward) Riccati recursions
- ☞ Require convexification in nonlinear programming

Remember that the IPM solution is always inaccurate!

IPM is NOT an active-set method

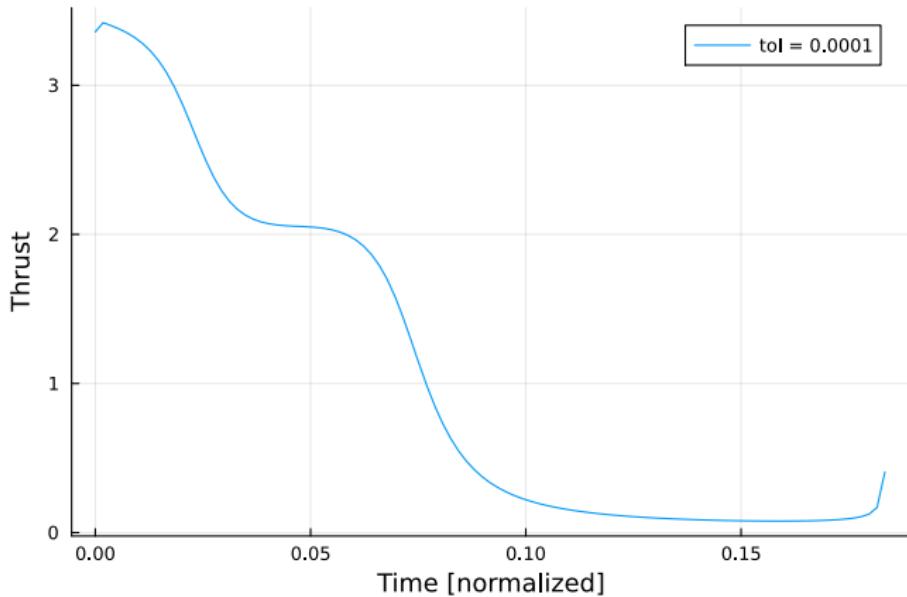


Figure: Optimal solution of the Goddard rocket problem with  $nh = 100$

Remember that the IPM solution is always inaccurate!

IPM is NOT an active-set method

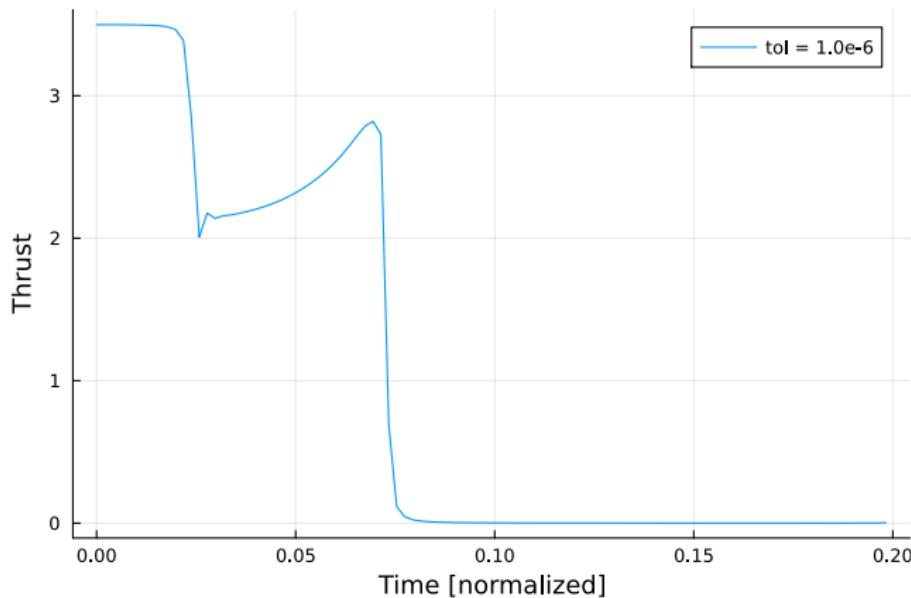


Figure: Optimal solution of the Goddard rocket problem with  $nh = 100$

Remember that the IPM solution is always inaccurate!

IPM is NOT an active-set method

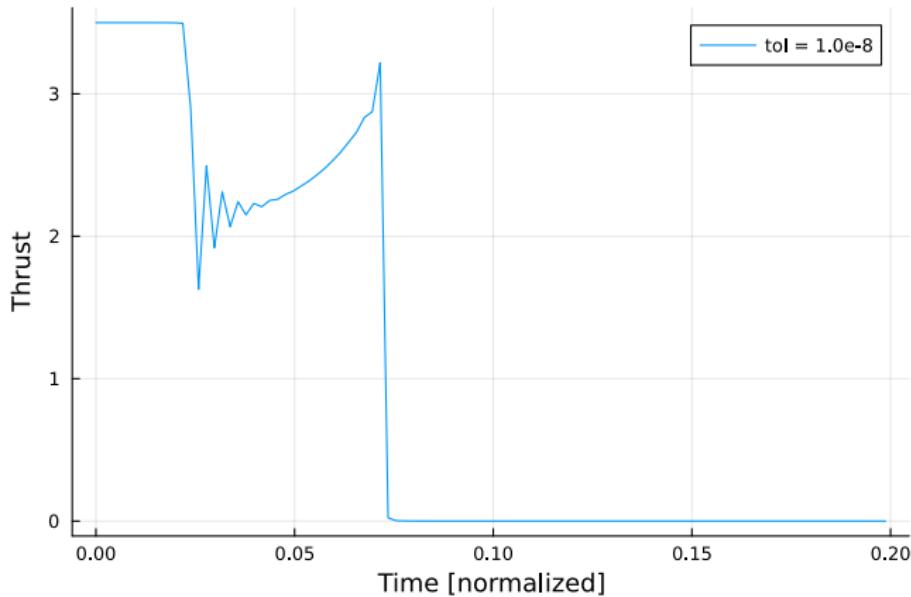


Figure: Optimal solution of the Goddard rocket problem with  $nh = 100$

Remember that the IPM solution is always inaccurate!

IPM is NOT an active-set method

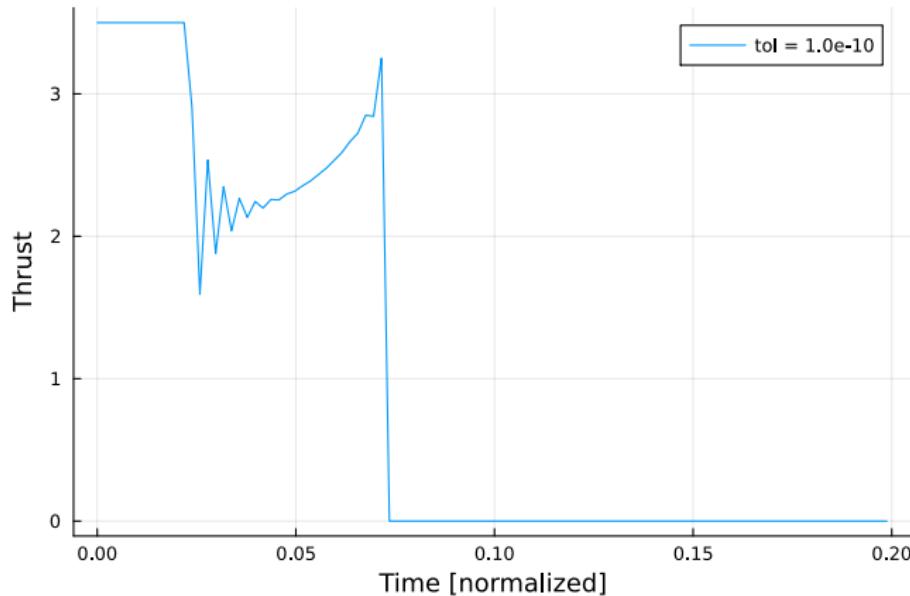


Figure: Optimal solution of the Goddard rocket problem with  $nh = 100$

Remember that the IPM solution is always inaccurate!

IPM is NOT an active-set method

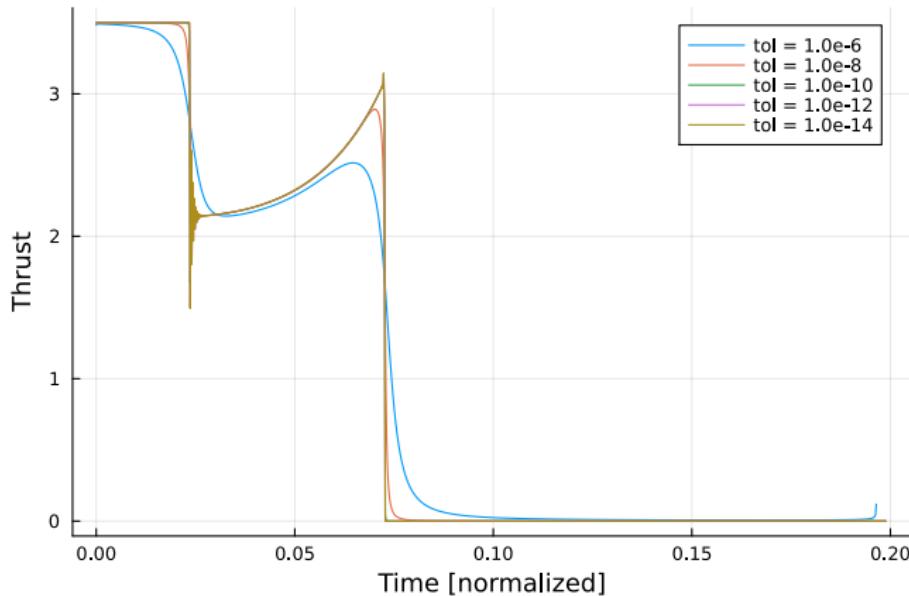


Figure: Optimal solution of the Goddard rocket problem with  $nh = 1000$

# Outline

What am I talking about when talking about optimal control?

Primal-dual interior-point method

How to solve the Newton systems on the GPU?

Using MadNLP and ExaModels on the GPU

# Why is it challenging?

## Current status

- ☞ Implementing a sparse linear solver on the GPU is highly non-trivial
- ☞ Before the release of NVIDIA cuDSS, there was no efficient sparse linear solver available on the GPU
- ☞ Solving the KKT systems in IPM with an iterative solver remains an open research question

Two solution methods for **GPU-accelerated** optimization method:

1. Rewrite the KKT system as a dense matrix ;
2. Use a pivoting-free factorization (at the price of sacrificing slightly the accuracy).

## Reduced-space approach: densify the solution!

Few degrees of freedom

### Proposition

Suppose the Jacobian  $\nabla g$  is full rank. Let  $Z$  be a matrix whose columns encode a basis of the Jacobian null-space ( $(\nabla g)Z = 0$ ). Then the condensed KKT system is equivalent to

$$Z^\top K Z d_u = \bar{r}_{red}$$

Oftentimes, the matrix  $Z^\top K Z$  is dense.

Example of null space for optimal control:

$$\nabla g = [G_x \ G_u] \implies Z = \begin{bmatrix} -G_x^{-1} G_u \\ I \end{bmatrix}$$

Works best if number of control is small (e.g. PDE constrained optimization)!

# Full-space approach: pivoting-free sparse linear solver

Many degrees of freedom

## How to avoid numerical pivoting?

Reformulate the KKT system either as

- ☞ a positive definite (PD) system

$$A \succ 0$$

- ☞ a symmetric quasidefinite (SQD) system

$$\begin{bmatrix} A & B^\top \\ B & -C \end{bmatrix} \quad \text{with} \quad A \succ 0, \ C \succ 0$$

In both cases, the sparse system can be factorized using only *static pivoting*, e.g. using a signed Cholesky factorization

$$A = PLD^\top P^\top$$

with  $P$  static pivoting,  $L$  lower triangular,  $D$  diagonal matrix.

# How to avoid numerical pivoting in nonlinear programming?

## Objective

How to avoid numerical pivoting in the algorithm?

We look again at the condensed KKT system (sparse symmetric indefinite):

$$\begin{bmatrix} \textcolor{teal}{K} & \nabla g^\top \\ \nabla g & 0 \end{bmatrix} \begin{bmatrix} d_w \\ d_y \end{bmatrix} = - \begin{bmatrix} r_1 \\ r_2 \end{bmatrix}$$

with the *condensed matrix*  $K = W + \nabla h^\top \Sigma_s \nabla h$ .

» Three strategies to avoid numerical pivoting:

1. LiftedKKT
2. HyKKT
3. NCL (Augmented Lagrangian)

## Strategy 1: LiftedKKT

Idea: equality relaxation

For a  $\tau > 0$  small enough, solve the relaxed problem

$$\min_{w \in \mathbb{R}^n} f(w) \quad \text{s.t.} \quad \begin{cases} -\tau \leq g(w) \leq \tau \\ h(w) \leq 0 \end{cases}$$

Reformulating the problem with slack variables:

$$\min_{w \in \mathbb{R}^n, s \in \mathbb{R}^m} f(w) \quad \text{subject to} \quad h^\tau(w) + s = 0, \quad s \geq 0$$

with  $h^\tau(w) = (g(w) - \tau, -g(w) - \tau, h(w))$

Condensed KKT system

The augmented KKT system is equivalent to

$$K_\tau d_w = -r_1 + (\nabla h^\tau)^\top (\Sigma_s r_4 + r_2)$$

with the *condensed matrix*  $K_\tau = W + (\nabla h^\tau)^\top \Sigma_s (\nabla h^\tau)$ .

→ the condensed KKT system can be solved without numerical pivoting!

## Strategy 2: HyKKT (aka Golub & Greif method)

Idea: augmented Lagrangian reformulation

For  $\gamma > 0$ , the condensed KKT system is equivalent to

$$\begin{bmatrix} \mathcal{K}_\gamma & \nabla g^\top \\ \nabla g & 0 \end{bmatrix} \begin{bmatrix} d_w \\ d_y \end{bmatrix} = - \begin{bmatrix} r_1 + \gamma \nabla g^\top r_2 \\ r_2 \end{bmatrix}$$

with  $\mathcal{K}_\gamma = K + \gamma \nabla g^\top \nabla g$

- ✓ For  $\gamma$  large-enough the matrix  $\mathcal{K}_\gamma$  is positive definite
  - ☛ solve the condensed KKT system using the normal equations:

$$(\nabla g) \mathcal{K}_\gamma^{-1} (\nabla g)^\top d_y = w_2 - \mathcal{K}_\gamma^{-1} (w_1 + \gamma \nabla g^\top w_2)$$

- ✓ Keep  $\mathcal{K}_\gamma^{-1}$  implicit by solving the normal equations *iteratively* with a conjugate gradient (CG) algorithm!
- ✓ For large  $\gamma$ , CG converges in few iterations

## Strategy 3: NCL

Augmented Lagrangian methods have always been popular in optimal control

At iteration  $k$ , the algorithm solves:

$$\begin{aligned} & \min_{w,r} f(w) - (y_k^e)^\top r + \frac{\rho_k}{2} \|r\|^2 \\ \text{subject to } & g(w) + r_g = 0, \\ & h(w) + r_h \leq 0, \end{aligned} \tag{NCL}_k$$

with  $r = (r_g, r_h)$  regularization variables

- Subproblem  $(NCL_k)$  is always feasible, solvable by IPM!
- Only the objective changes between  $k$  and  $k + 1$
- Regularization  $r$  stabilizes internal IPM iterations

### NCL algorithm $\equiv$ Auglag algorithm

- Solve  $(NCL_k)$  down to a tolerance  $\omega_k$
- Update parameters as
  - If  $\|r_{k+1}\| \leq \eta_k$ , set  $y_{k+1}^e = y_k^e - \rho_k r_{k+1}$
  - Else  $\rho_{k+1} = 10 \times \rho_k$ .

Algorithm NCL is more robust, but converges in more iterations than IPM.

## Strategy 3: Stabilized KKT system

Rockafellar: Augmented Lagrangian adds a natural dual regularization to the KKT system!

For the two diagonal matrices  $C_g := \frac{1}{\rho_k} I$  and  $C_h := \frac{1}{\rho_k} I + Z^{-1}S$ , let

$$\begin{bmatrix} W & \nabla g^\top & \nabla h^\top \\ \nabla g & -C_g & 0 \\ \nabla h & 0 & -C_h \end{bmatrix} \begin{bmatrix} d_w \\ d_y \\ d_z \end{bmatrix} = - \begin{bmatrix} r_1 \\ r_2 \\ r_3 \end{bmatrix} \quad (K_{2r})$$

### Observation

NCL can compute its descent direction by solving  $K_{2r}$ .

- ☞ For a penalty  $\rho_k$  high-enough, there exists a LDL factorization for  $K_{2r}$
- ☞ Otherwise, use a *pivot regularization* strategy inside LDL to return the factorization of a *perturbed matrix*
- ☞ Recover the original descent direction using iterative refinement

# Outline

What am I talking about when talking about optimal control?

Primal-dual interior-point method

How to solve the Newton systems on the GPU?

Using MadNLP and ExaModels on the GPU

# GPU-accelerated optimization

New solvers are blossoming everywhere

Most solvers are specialized on convex problems:

☞ First-order

- cuPDLP and all affiliated variants
- OSQP
- cuHALLaR / cuLORADS

(LP, QP)  
(QP)  
(low-rank SDP)

☞ Second-order

- QPTH
- CuClarabel

(dense QP)  
(conic)

MadSuite: an optimization software suite for GPUs

[madsuite.org](http://madsuite.org)

- **ExaModels** (NLP modeler)
- **MadNLP** (NLP)
- **MadNCL** (degenerate NLP)
- **MadIPM** (LP)

# State-of-the-art solvers for optimal control

Support for GPU-accelerated optimization is coming

## Observations

State-of-the-art solvers are *all* exploiting the problem's dynamic structure

- ☞ Differential dynamic programming (Altro, ProxDDP)
- ☞ Condensation strategy (HPIPM)
- ☞ Riccati recursion (FATROP)

Further, the ecosystem leverages mature tools:

- BLAS library for embedded system: BLASFEO
- Efficient modeler: Casadi

## Question

So, what do we gain by solving optimal control instances on the GPU?

- Large-scale optimization problems **almost always have repetitive patterns**

$$\min_{x^b \leq x \leq x^u} \sum_{l \in [L]} \sum_{i \in [I_l]} f^{(l)}(x; p_i^{(l)}) \quad (\text{SIMD abstraction})$$

$$\text{subject to } \left[ g^{(m)}(x; q_j) \right]_{j \in [J_m]} + \sum_{n \in [N_m]} \sum_{k \in [K_n]} h^{(n)}(x; s_k^{(n)}) = 0, \quad \forall m \in [M]$$

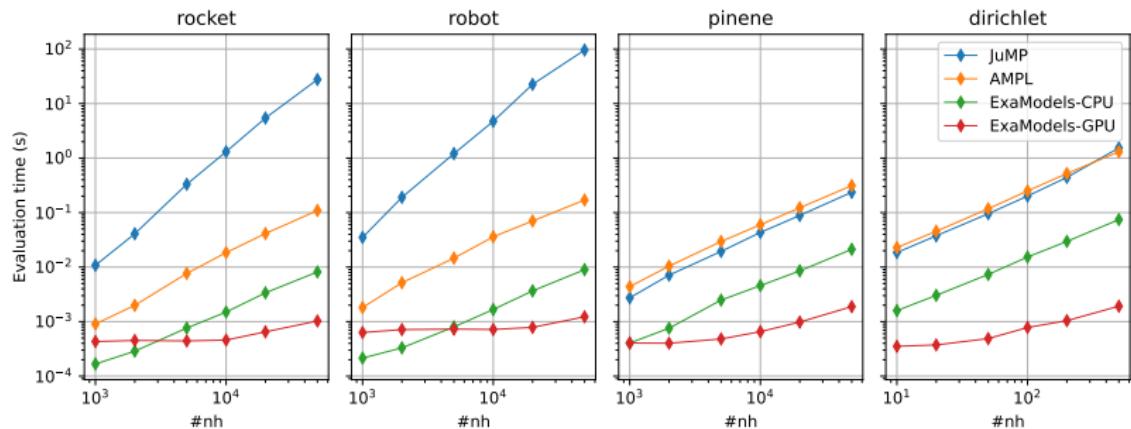
- Repeated patterns are made available by specifying the models as **iterable objects**

```
constraint(c, 3 * x[i+1]^3 + 2 * sin(x[i+2])) for i = 1:N-2)
```

- For each repetitive pattern, the derivative evaluation kernel is constructed & compiled, and **executed in parallel over multiple data**

☞ **Perfect settings for optimal control**  
**(repeated pattern: cost and dynamics)**

# How fast can we get with ExaModels?



**Figure:** Time to evaluate the derivatives (Jacobian + Hessian) on various COPS instances with a dynamic structure

## Now, combining ExaModels with NVIDIA cuDSS

N.B.: here, we are not exploiting the problem's dynamic structure

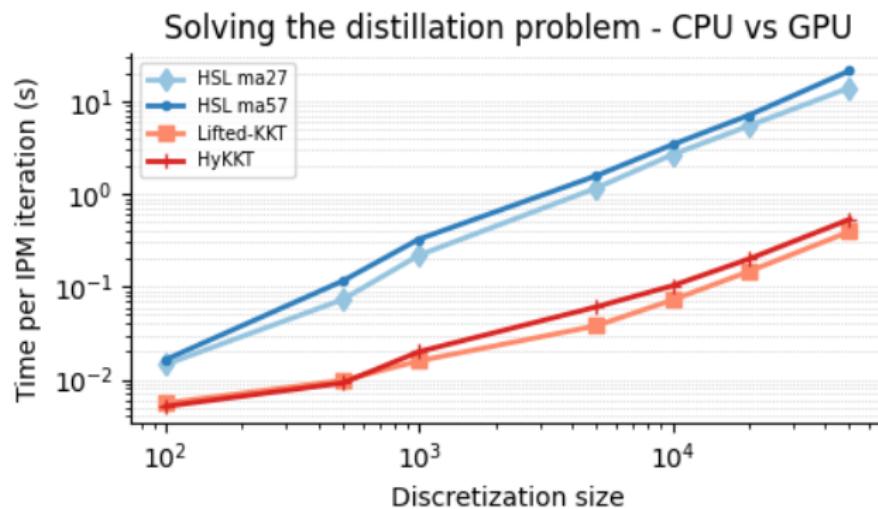


Figure: Solving the infamous distillation column instance

# How expensive should be your GPU?

## Observation

- No need to buy a professional GPU to get fast performance with ExaModels+cuDSS

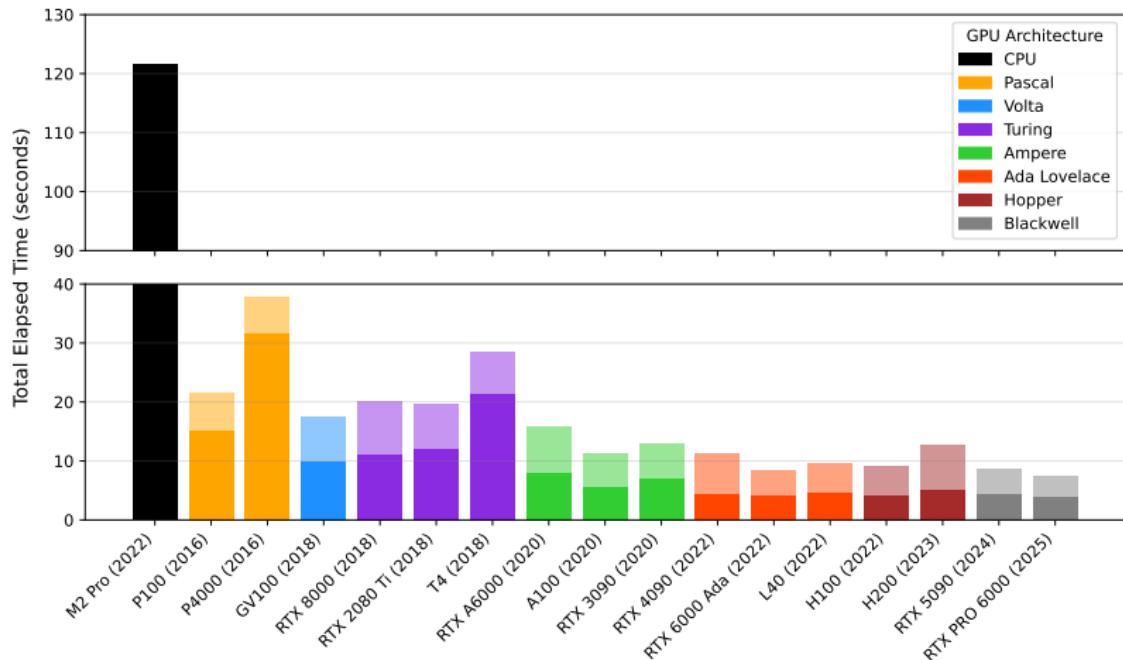


Figure: Time to optimality, here for a large-scale optimal power flow instance.

## Next step: Solving the system in parallel

Current research effort

The KKT system is a block tridiagonal system:

$$\begin{bmatrix} A_{1,1} & A_{2,1}^\top \\ A_{2,1}^\top & \ddots & \ddots \\ \ddots & \ddots & A_{N,N-1}^\top \\ & & A_{N,N-1} & A_{N,N} \end{bmatrix} \begin{bmatrix} d_{w1} \\ \vdots \\ \vdots \\ d_{wN} \end{bmatrix} = \begin{bmatrix} r_1 \\ \vdots \\ \vdots \\ r_N \end{bmatrix}$$

Two different path way:

- ☞ Parallel Cycling Reduction (PCR)
- ☞ Partitionned Dynamic Programming

Integration is under way in MadNLP!

# Conclusion

GPU-accelerated optimal control is currently happening!

## ExaModels

Fast evaluations of derivatives in nonlinear models

[frapac.github.io/tutorials/powertech/](https://frapac.github.io/tutorials/powertech/)

## MadNLP

Fast solution of nonlinear programs on the GPU