

# MadIPM: a GPU-accelerated solver for linear programming

*Why second-order methods remain relevant*

**François Pacaud**

CAS, Mines Paris - PSL

*PGMO Days 2025*

# Who are we?

<https://madsuite.org/>



- Alexis Montoisin @ Argonne National Laboratory
- François Pacaud @ MINES Paris-PSL
- Sungho Shin @ MIT
- Mihai Anitescu @ Argonne National Laboratory

And friends: Dominique Orban and JSO

# What do we want to discuss today?

## Observations:

- GPU-accelerated optimization is gaining momentum
- Most recent developments are using first-order methods (ADMM, PDLP)
- New generation of sparse linear solvers on the GPU (cuDSS)

## Research question

Are second-order methods effective at solving large-scale LPs on the GPU?

## Newest GPU-accelerated solvers for LPs:

- cuPDLP
- NVIDIA cuOPT
- cuClarabel
- MadIPM!

## MadIPM — GPU example using JuMP

Available on github to solve your own LPs!

[github.com/MadNLP/MadIPM.jl](https://github.com/MadNLP/MadIPM.jl)

```
using JuMP, MadIPM
using CUDA, KernelAbstractions, MadNLPGPU

c = rand(10)
model = Model(MadIPM.Optimizer)

# GPU settings
set_optimizer_attribute(model, "array_type", CuVector{Float64})
set_optimizer_attribute(model, "linear_solver", MadNLPGPU.
    CUDSSSolver)

@variable(model, 0 <= x[1:10], start=0.5)
@constraint(model, sum(x) == 1.0)
@objective(model, Min, c' * x)

JuMP.optimize!(model)
```

## How to solve a LP with interior-point?

We define a LP in **standard form** as

$$\min_{x \in \mathbb{R}^n} c^\top x \quad \text{subject to} \quad Ax = b, \quad x \geq 0, \quad (\text{LP})$$

with  $c \in \mathbb{R}^n$ ,  $A \in \mathbb{R}^{m \times n}$ ,  $b \in \mathbb{R}^m$  problem's data

Any interior-point solver uses the three following items to solve **(LP)**:

1. A good **pre-processing**
2. A Newton method to track the **central path**
3. A sparse linear solver to compute the **Newton direction**

The **Mehrotra's predictor-corrector** is the workhorse IPM method for **(LP)**

## Step 1: pre-processing

We reformulate (LP) in a suitable form for IPM (this step is **very important** to get a reliable convergence)

1. Call a pre-processing routine to remove fixed variables and dependent constraints
2. Scale the problem's data by applying the **Ruiz equilibration** method on the constraint matrix  $A$
3. If appropriate, reformulate the problem in standard form
4. Find a good initial point using Mehrotra's heuristic  
*(push away from boundary to get long steps in the first iterations)*

## KKT equations

We write the KKT equations for the LP in standard form

The primal-dual point  $w := (x, y, z)$  is a solution of (LP) if and only if

$$A^\top y - z = 0$$

$$Ax - b = 0$$

$$0 \leq x \perp z \geq 0$$

For a barrier parameter  $\mu > 0$ , we say that  $w$  is on the central path if  $(x, z) > 0$  and

$$A^\top y - z = 0$$

$$Ax - b = 0$$

$$XZe = \mu e$$

### Primal-dual interior-point method

≡ Track the central path using Newton method

## Step 2: Mehrotra predictor-corrector

We have to ensure that we track the central path

For a given primal-dual iterate  $w = (x, y, z)$ , define the **current barrier parameter** (average complementarity) as:

$$\mu = \frac{z^\top x}{n}.$$

**Affine step:** Compute  $\Delta^{\text{aff}}$  by solving

$$\begin{bmatrix} 0 & A^\top & -I \\ A & 0 & 0 \\ Z & 0 & X \end{bmatrix} \begin{bmatrix} \Delta x^{\text{aff}} \\ \Delta y^{\text{aff}} \\ \Delta z^{\text{aff}} \end{bmatrix} = - \begin{bmatrix} c + A^\top y - z \\ Ax - b \\ Xze \end{bmatrix}.$$

**Corrector step:** Compute  $\Delta^{\text{corr}}$  by solving

$$\begin{bmatrix} 0 & A^\top & -I \\ A & 0 & 0 \\ Z & 0 & X \end{bmatrix} \begin{bmatrix} \Delta x^{\text{corr}} \\ \Delta y^{\text{corr}} \\ \Delta z^{\text{corr}} \end{bmatrix} = - \begin{bmatrix} 0 \\ 0 \\ \sigma\mu e - \Delta Z^{\text{aff}} \Delta X^{\text{aff}} e \end{bmatrix},$$

with  $\sigma$  given by a heuristic

**Update:** For  $\Delta^k = \Delta^{\text{aff}} + \Delta^{\text{corr}}$ , set

$$w^{k+1} = w^k + \alpha^k \Delta^k$$

with  $\alpha$  a step computed using a fraction-to-boundary rule

## Step 3: sparse linear solver

The affine step and the corrector step are both solving the **unsymmetric linear system**:

$$\begin{bmatrix} 0 & A^\top - I \\ A & 0 \\ Z & X \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \\ \Delta z \end{bmatrix} = \begin{bmatrix} r_1 \\ r_2 \\ r_3 \end{bmatrix}.$$

### Augmented KKT system

The unsymmetric system reduces to:

$$\begin{bmatrix} \Sigma & A^\top \\ A & 0 \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} = \begin{bmatrix} r_1 + X^{-1}r_3 \\ r_2 \end{bmatrix},$$

with the diagonal matrix  $\Sigma := X^{-1}Z$ .

### Normal KKT system

We can also eliminate  $\Delta y$  to recover the positive-definite system:

$$(A\Sigma^{-1}A^\top)\Delta y = A\Sigma^{-1}(r_1 + X^{-1}r_3) - r_2.$$

## Primal-dual regularization

It is the key for GPU performance!

Both the augmented and normal KKT systems have their own issues if:

- Free variables
- Rank-deficient Jacobian  $A$
- Dense rows in  $A$  leads to a dense matrix  $A\Sigma^{-1}A^\top$

Instead, we regularize the system using two small positive parameters  $(\rho, \delta) > 0$ .

The system becomes:

$$\begin{bmatrix} \Sigma + \rho I & A^\top \\ A & -\delta I \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix} = \begin{bmatrix} r_1 + X^{-1}r_3 \\ r_2 \end{bmatrix}.$$

The matrix is **symmetric quasi-definite** (SQD), meaning that it is strongly factorizable using a signed Cholesky factorization.

### Proposition

This is equivalent to solve the **regularized LP**:

$$\min_{x,r} c^\top x + \frac{\rho}{2} \|x\|^2 + \frac{1}{2\delta} \|r\|^2$$

$$\text{subject to } Ax + r = b, \quad x \geq 0$$

## MadIPM: performance profile

### Benchmark MadIPM




Figure: Benchmarking MadIPM, Gurobi and HiGHS on 174 large-scale LP instances from MIPLIB

*The upper the better*

## MadIPM: raw performance




Figure: Zoom on the largest instances

*The lower the better*

# How expensive should be your GPU?

Image courtesy of Sungho Shin




Figure: Time to optimality, here for a large-scale optimal power flow instance.

**No need to buy a professional GPU for fast performance!**

# Conclusion

If you are interested, you can read the full article:

*Montoison, Pacaud, Shin, Anitescu.  
GPU Implementation of Second-Order Linear and Nonlinear Programming Solvers.  
2025.*

## Next step

Solve LPs in **batch** on the GPU!

- Solve multiple LPs **simultaneously**, assuming the same KKT sparsity pattern
- Reuse **symbolic analysis** across all sparse linear systems
- Different central paths → real-time rebalancing when some systems converge earlier.

## The one-billion dollars question

- How to solve large-scale MILPs on the GPU?